

CYCLE – 2

I Build a small Convolutional Neural Network (CNN) model using any of deep libraries for:

- a) Image Recognition/ Classification
- b) Digit Identification

Code:

a)

```
import tensorflow as tf
from tensorflow.keras import layers, models

# Load CIFAR-10 dataset
(train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.cifar10.load_data()

# Normalize pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0

# Define the CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
model.fit(train_images, train_labels, epochs=10, batch_size=64, validation_data=(test_images,
test_labels))
```

b)

```
import tensorflow as tf
from tensorflow.keras import layers, models
```

```

# Load MNIST dataset
mnist = tf.keras.datasets.mnist
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

# Normalize pixel values to be between 0 and 1
train_images, test_images = train_images / 255.0, test_images / 255.0

# Define the CNN model for digit identification
model = models.Sequential([
layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
layers.MaxPooling2D((2, 2)),
layers.Conv2D(64, (3, 3), activation='relu'),
layers.MaxPooling2D((2, 2)),
layers.Conv2D(64, (3, 3), activation='relu'),
layers.Flatten(),
layers.Dense(64, activation='relu'),
layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
loss='sparse_categorical_crossentropy',
metrics=['accuracy'])

# Reshape the input data to fit the CNN input shape
train_images = train_images.reshape(train_images.shape[0], 28, 28, 1)
test_images = test_images.reshape(test_images.shape[0], 28, 28, 1)

# Train the model
model.fit(train_images, train_labels, epochs=5, batch_size=64, validation_data=(test_images,
test_labels))

```

Output:

a)

Epoch 6/10

782/782 [=====] - 4s 5ms/step - loss: 0.8482 - accuracy: 0.7076 -
val_loss: 0.9099 - val_accuracy: 0.6842

Epoch 7/10

782/782 [=====] - 4s 5ms/step - loss: 0.7968 - accuracy: 0.7231 -
val_loss: 0.8801 - val_accuracy: 0.6972

Epoch 8/10

782/782 [=====] - 4s 6ms/step - loss: 0.7487 - accuracy: 0.7395 -
val_loss: 0.8674 - val_accuracy: 0.6994

Epoch 9/10

```

782/782 [=====] - 4s 5ms/step - loss: 0.7155 - accuracy: 0.7491 -
val_loss: 0.8860 - val_accuracy: 0.6953
Epoch 10/10
782/782 [=====] - 4s 5ms/step - loss: 0.6783 - accuracy: 0.7652 -
val_loss: 0.8507 - val_accuracy: 0.7162
<keras.src.callbacks.History at 0x7e817c7c0b50>

```

b)

```

Epoch 1/5 938/938 [=====] - 7s 5ms/step - loss: 0.1821
- accuracy: 0.9449 - val_loss: 0.0657 - val_accuracy: 0.9812 Epoch 2/5 938/938
[=====] - 5s 5ms/step - loss: 0.0527 - accuracy:
0.9838 - val_loss: 0.0462 - val_accuracy: 0.9854 Epoch 3/5 938/938
[=====] - 5s 5ms/step - loss: 0.0392 - accuracy:
0.9878 - val_loss: 0.0293 - val_accuracy: 0.9895 Epoch 4/5 938/938
[=====] - 5s 5ms/step - loss: 0.0287 - accuracy:
0.9906 - val_loss: 0.0390 - val_accuracy: 0.9862 Epoch 5/5 938/938
[=====] - 5s 6ms/step - loss: 0.0237 - accuracy:
0.9924 - val_loss: 0.0313 - val_accuracy: 0.9891

```

II How to use Pre-trained CNN models for feature extraction.

Code:

```

import numpy as np
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.resnet50 import preprocess_input
from tensorflow.keras.models import Model
import matplotlib.pyplot as plt

# Load pre-trained ResNet50 model
resnet50_model = ResNet50(weights='imagenet', include_top=False)

# Define a new model with ResNet50 base
model = Model(inputs=resnet50_model.input,
outputs=resnet50_model.get_layer('conv5_block3_out').output)

# Function to extract features from images using ResNet50
def extract_features(img_path):
    img = image.load_img(img_path, target_size=(224, 224)) # ResNet50 expects images of size 224x224
    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)
    x = preprocess_input(x) # Preprocess the input data to align with the way ResNet50 was trained
    features = model.predict(x)

```

```
return features.flatten() # Flatten the features to use them as input to another classifier
```

```
# Example usage
```

```
img_path = '/content/drive/MyDrive/cat.png'
```

```
features = extract_features(img_path)
```

```
print("Shape of extracted features:", features.shape)
```

```
# Function to plot histogram of extracted features
```

```
def plot_features_histogram(features):
```

```
plt.hist(features, bins=50)
```

```
plt.title('Histogram of Extracted Features')
```

```
plt.xlabel('Feature Value')
```

```
plt.ylabel('Frequency')
```

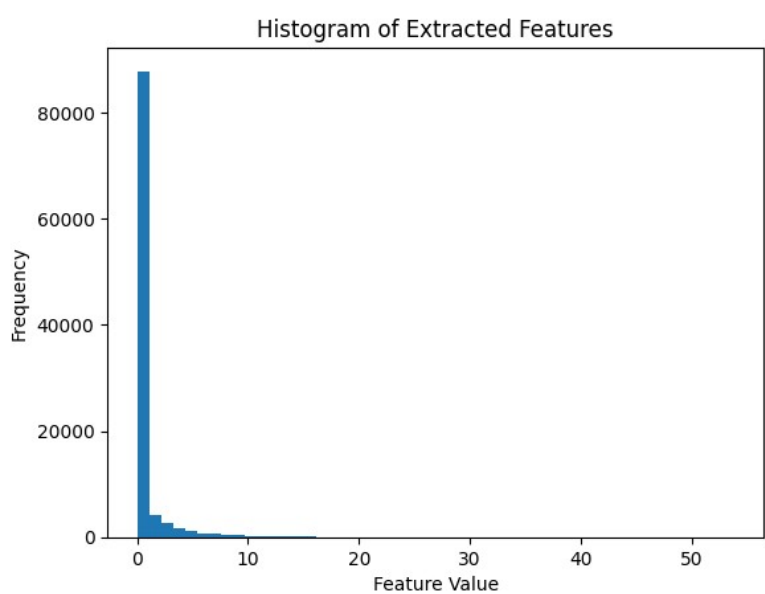
```
plt.show()
```

```
# Example usage
```

```
plot_features_histogram(features)
```

Output:

```
1/1 [=====] - 2s 2s/step Shape of extracted features:  
(100352, )
```



III Implementation of Pre-trained CNN models using transfer learning for classification/object detections.

a) AlexNet

b) VGG-16

Code:

a)

```
import torch

import torch.nn as nn
import torch.optim as optim
import torchvision.transforms as transforms
import torchvision.datasets as datasets
import torchvision.models as models
from torch.utils.data import DataLoader

# Set device (GPU if available, else CPU)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

# Define transformation for data augmentation and normalization
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

# Load CIFAR-10 dataset
train_dataset = datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
test_dataset = datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)

# Create data loaders
train_loader = DataLoader(dataset=train_dataset, batch_size=64, shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=64, shuffle=False)

# Load pre-trained AlexNet model
alexnet_model = models.alexnet(pretrained=True)

# Freeze parameters of pre-trained layers
for param in alexnet_model.parameters():
    param.requires_grad = False

# Modify the last fully connected layer for CIFAR-10 classification (10 classes)
num_features = alexnet_model.classifier[6].in_features
alexnet_model.classifier[6] = nn.Linear(num_features, 10)
```

```

# Move model to device
alexnet_model = alexnet_model.to(device)

# Define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(alexnet_model.parameters(), lr=0.001)

# Train the model
num_epochs = 5
for epoch in range(num_epochs):
    running_loss = 0.0
    for i, (inputs, labels) in enumerate(train_loader):
        inputs = inputs.to(device)
        labels = labels.to(device)

        optimizer.zero_grad()

        outputs = alexnet_model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()

    if (i+1) % 100 == 0:
        print(f'Epoch [{epoch+1}/{num_epochs}], Step [{i+1}/{len(train_loader)}], Loss:
        {running_loss/100}')
        running_loss = 0.0

# Evaluate the model
alexnet_model.eval()
correct = 0
total = 0
with torch.no_grad():
    for inputs, labels in test_loader:
        inputs = inputs.to(device)
        labels = labels.to(device)
        outputs = alexnet_model(inputs)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total
print(f'Accuracy of the network on the 10000 test images: {accuracy}%)

```

b)

```
import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.applications import VGG16
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam

# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

# Load pre-trained VGG16 model (without top fully connected layers)
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))

# Freeze the convolutional layers
for layer in base_model.layers:
    layer.trainable = False

# Add custom classification layers
x = GlobalAveragePooling2D()(base_model.output)
x = Dense(256, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

# Create the model
model = Model(inputs=base_model.input, outputs=predictions)

# Compile the model
model.compile(optimizer=Adam(), loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Train the model
model.fit(x_train, y_train, epochs=50, batch_size=64, validation_data=(x_test, y_test))
```

Output:

a)

```
Epoch [5/5], Step [100/782], Loss: 0.5746017411351204
Epoch [5/5], Step [200/782], Loss: 0.5881411558389664
Epoch [5/5], Step [300/782], Loss: 0.6102184489369392
Epoch [5/5], Step [400/782], Loss: 0.5950301320850849
Epoch [5/5], Step [500/782], Loss: 0.6262218597531318
Epoch [5/5], Step [600/782], Loss: 0.5968183997273445
```

Epoch [5/5], Step [700/782], Loss: 0.6252494990825653
Accuracy of the network on the 10000 test images: 82.13%

b)

Epoch 46/50

782/782 [=====] - 11s 14ms/step - loss: 0.3186 -
accuracy: 0.8953 - val_loss: 1.7425 - val_accuracy: 0.5968

Epoch 47/50

782/782 [=====] - 10s 13ms/step - loss: 0.3130 -
accuracy: 0.8953 - val_loss: 1.7340 - val_accuracy: 0.6060

Epoch 48/50

782/782 [=====] - 10s 13ms/step - loss: 0.3026 -
accuracy: 0.8994 - val_loss: 1.7971 - val_accuracy: 0.5943

Epoch 49/50

782/782 [=====] - 11s 14ms/step - loss: 0.2939 -
accuracy: 0.9028 - val_loss: 1.7883 - val_accuracy: 0.6012

Epoch 50/50

782/782 [=====] - 11s 14ms/step - loss: 0.2841 -
accuracy: 0.9060 - val_loss: 1.8380 - val_accuracy: 0.6009

IV Practicing various strategies of fine tuning.

Code:

```
import tensorflow as tf

from tensorflow.keras.datasets import cifar10
from tensorflow.keras.applications import VGG16
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ModelCheckpoint

# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

# Load pre-trained VGG16 model
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))
```



```

# Freeze layers except the last convolutional block
for layer in base_model.layers[:-4]:
    layer.trainable = False

# Add custom classification layers
x = GlobalAveragePooling2D()(base_model.output)
x = Dense(256, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

# Create the model
model = Model(inputs=base_model.input, outputs=predictions)

# Compile the model
model.compile(optimizer=Adam(lr=1e-4), loss='sparse_categorical_crossentropy',
metrics=['accuracy'])

# Print model summary
model.summary()

# Train the model
history = model.fit(x_train, y_train, epochs=10, batch_size=64, validation_data=(x_test, y_test))

# Evaluate the model
test_loss, test_acc = model.evaluate(x_test, y_test)
print("Test Accuracy:", test_acc)

```

Output:

```

Total params: 14848586 (56.64 MB)
Trainable params: 7213322 (27.52 MB)
Non-trainable params: 7635264 (29.13 MB)

```

Epoch 6/10

```

782/782 [=====] - 16s 20ms/step - loss: 0.4759 -
accuracy: 0.8338 - val_loss: 0.8067 - val_accuracy: 0.7454

```

Epoch 7/10

```

782/782 [=====] - 17s 22ms/step - loss: 0.4210 -
accuracy: 0.8527 - val_loss: 0.8570 - val_accuracy: 0.7322

```

Epoch 8/10

```

782/782 [=====] - 16s 20ms/step - loss: 0.3660 -
accuracy: 0.8697 - val_loss: 0.9557 - val_accuracy: 0.7352

```

Epoch 9/10

782/782 [=====] - 16s 21ms/step - loss: 0.3181 - accuracy: 0.8870 - val_loss: 0.9987 - val_accuracy: 0.7391

Epoch 10/10

782/782 [=====] - 16s 20ms/step - loss: 0.2770 - accuracy: 0.9025 - val_loss: 1.0109 - val_accuracy: 0.7410 313/313

[=====] - 3s 9ms/step - loss: 1.0109 - accuracy: 0.7410

V Implementing Generative Models:

- a) Autoencoder for image reconstruction
- b) Word Prediction using RNN
- c) Image Captioning

Code:

a)

```
import numpy as np

import matplotlib.pyplot as plt
from tensorflow.keras.layers import Input, Dense, Conv2D, MaxPooling2D, UpSampling2D
from tensorflow.keras.models import Model
from tensorflow.keras.datasets import mnist

# Load MNIST dataset
(x_train, _), (x_test, _) = mnist.load_data()

# Normalize pixel values to be between 0 and 1
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = np.reshape(x_train, (len(x_train), 28, 28, 1))
x_test = np.reshape(x_test, (len(x_test), 28, 28, 1))

# Define the autoencoder model
input_img = Input(shape=(28, 28, 1))
x = Conv2D(16, (3, 3), activation='relu', padding='same')(input_img)
x = MaxPooling2D((2, 2), padding='same')(x)
x = Conv2D(8, (3, 3), activation='relu', padding='same')(x)
x = MaxPooling2D((2, 2), padding='same')(x)
x = Conv2D(8, (3, 3), activation='relu', padding='same')(x)
```

```

encoded = MaxPooling2D((2, 2), padding='same')(x)

# Decoder
x = Conv2D(8, (3, 3), activation='relu', padding='same')(encoded)
x = UpSampling2D((2, 2))(x)
x = Conv2D(8, (3, 3), activation='relu', padding='same')(x)
x = UpSampling2D((2, 2))(x)
x = Conv2D(16, (3, 3), activation='relu')(x)
x = UpSampling2D((2, 2))(x)
decoded = Conv2D(1, (3, 3), activation='sigmoid', padding='same')(x)

# Create autoencoder model
autoencoder = Model(input_img, decoded)
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# Train the model
autoencoder.fit(x_train, x_train,
epochs=10,
batch_size=128,
shuffle=True,
validation_data=(x_test, x_test))

# Generate reconstructed images
decoded_imgs = autoencoder.predict(x_test)

# Plot some of the reconstructed images
n = 10 # Number of images to display
plt.figure(figsize=(20, 4))
for i in range(n):
# Display original images
ax = plt.subplot(2, n, i + 1)
plt.imshow(x_test[i].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)

# Display reconstructed images
ax = plt.subplot(2, n, i + 1 + n)
plt.imshow(decoded_imgs[i].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()

```

b)

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
import numpy as np
import random
import string
import re
from sklearn.model_selection import train_test_split

# Sample list of sentences
text_data = [
    "I like to eat apples and oranges",
    "Apples and oranges are delicious fruits",
    "I prefer oranges over apples",
    "Oranges are juicy and delicious"
]

# Preprocess and tokenize text data
def preprocess_text(text):
    text = text.lower()
    text = re.sub(r'\d+', '', text) # Remove numbers
    text = text.translate(str.maketrans("", "", string.punctuation)) # Remove punctuation
    return text.split()

# Preprocess and tokenize text data
tokens = []
for sentence in text_data:
    tokens.extend(preprocess_text(sentence))

# Create sequences of input and target pairs
seq_length = 5
sequences = []
for i in range(len(tokens) - seq_length):
    sequences.append((tokens[i:i+seq_length], tokens[i+seq_length]))

# Create word-to-index and index-to-word mappings
word_to_idx = {word: idx for idx, word in enumerate(set(tokens))}
idx_to_word = {idx: word for word, idx in word_to_idx.items()}
vocab_size = len(word_to_idx)

# Define a PyTorch Dataset
class TextDataset(Dataset):
    def __init__(self, sequences):
        self.sequences = sequences
```

```

def __len__(self):
    return len(self.sequences)

def __getitem__(self, idx):
    input_seq, target = self.sequences[idx]
    input_idx = torch.tensor([word_to_idx[word] for word in input_seq], dtype=torch.long)
    target_idx = torch.tensor(word_to_idx[target], dtype=torch.long)
    return input_idx, target_idx

# Split data into train and test sets
train_data, test_data = train_test_split(sequences, test_size=0.2, random_state=42)

# Create DataLoader instances
train_loader = DataLoader(TextDataset(train_data), batch_size=64, shuffle=True)
test_loader = DataLoader(TextDataset(test_data), batch_size=64, shuffle=False)

# Define the LSTM model
class WordPredictionModel(nn.Module):
    def __init__(self, vocab_size, embedding_dim, hidden_dim):
        super(WordPredictionModel, self).__init__()
        self.embedding = nn.Embedding(vocab_size, embedding_dim)
        self.lstm = nn.LSTM(embedding_dim, hidden_dim, batch_first=True)
        self.fc = nn.Linear(hidden_dim, vocab_size)

    def forward(self, x):
        embedded = self.embedding(x)
        lstm_out, _ = self.lstm(embedded)
        out = self.fc(lstm_out[:, -1, :])
        return out

# Initialize the model, loss function, and optimizer
embedding_dim = 100
hidden_dim = 128
model = WordPredictionModel(vocab_size, embedding_dim, hidden_dim)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training the model
num_epochs = 10
for epoch in range(num_epochs):
    running_loss = 0.0
    for input_idx, target_idx in train_loader:
        optimizer.zero_grad()
        output = model(input_idx)
        loss = criterion(output, target_idx)
        loss.backward()
        optimizer.step()

```

```

running_loss += loss.item()
print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {running_loss / len(train_loader):.4f}')

# Function to generate next words
def generate_next_words(seed_text, next_words):
    with torch.no_grad():
        seed_seq = seed_text.split()
        for _ in range(next_words):
            input_idx = torch.tensor([word_to_idx[word] for word in seed_seq], dtype=torch.long).unsqueeze(0)
            output = model(input_idx)
            predicted_idx = torch.argmax(output, dim=1).item()
            next_word = idx_to_word.get(predicted_idx, '<UNK>')
            seed_seq.append(next_word)
        return ' '.join(seed_seq)

# Generate and print next words
seed_text = "apples are delicious"
print(generate_next_words(seed_text, 5))

```

c)

```

import matplotlib.pyplot as plt

import torch
from torchvision.transforms import transforms
from PIL import Image
from transformers import BlipProcessor, BlipForConditionalGeneration
import nltk
nltk.download('punkt')

# Load the pre-trained image captioning model
processor = BlipProcessor.from_pretrained("Salesforce/blip-image-captioning-base")
model = BlipForConditionalGeneration.from_pretrained("Salesforce/blip-image-captioning-base")

# Load and preprocess the image
image_path = "/content/cat.jfif"
image = Image.open(image_path).convert("RGB") # Convert to RGB
preprocess = transforms.Compose([
    transforms.Resize((256, 256)),
    transforms.ToTensor(),
])
input_tensor = preprocess(image).unsqueeze(0)

# Generate captions
with torch.no_grad():

```

```

captions = model.generate(pixel_values=input_tensor)

# Decode the generated captions
caption_text = processor.decode(captions[0], skip_special_tokens=True)

# Print the generated caption
print("Generated Caption:", caption_text)

plt.imshow(image)
plt.axis('off')
plt.show

```


Output:

a)

```

Epoch 6/10
469/469 [=====] - 4s 8ms/step - loss: 0.1126 -
val_loss: 0.1102
Epoch 7/10
469/469 [=====] - 3s 6ms/step - loss: 0.1108 -
val_loss: 0.1085
Epoch 8/10
469/469 [=====] - 3s 6ms/step - loss: 0.1091 -
val_loss: 0.1070
Epoch 9/10
469/469 [=====] - 3s 6ms/step - loss: 0.1079 -
val_loss: 0.1058
Epoch 10/10
469/469 [=====] - 4s 8ms/step - loss: 0.1068 -
val_loss: 0.1052
313/313 [=====] - 1s 2ms/step

```



b)

Epoch [6/10], Loss: 2.2166

Epoch [7/10], Loss: 2.1457

Epoch [8/10], Loss: 2.0729

Epoch [9/10], Loss: 1.9981

Epoch [10/10], Loss: 1.9212

apples are delicious fruits fruits prefer oranges apples

c)

Generated Caption: a group of cats sitting in a row

